

# Does Health Spending Reduce Child Mortality?

A causal analysis of World Bank / Our World in Data health-expenditure and child-mortality indicators, produced end-to-end with the aitiolin Causality Agent.

**Dataset:** worldbank\_health\_mortality\_2019.csv (8 raw variables, 186 countries, single cross-section, year 2019)

**Treatment:** health\_expenditure\_per\_capita (current int-\$, PPP)

**Outcome:** child\_mortality\_rate (deaths per 1,000 live births, under age 5)

**Confounders identified:** gdp\_per\_capita, region

**Model source:** LLM-assisted (GPT) causal-model draft, statistically verified against the data

## Headline result

Once GDP per capita and world region are adjusted for, the estimated effect of health spending on child mortality is **small and statistically indistinguishable from zero**: raising health expenditure by \$1,000 per capita is associated with roughly **-0.015 fewer deaths per 1,000 live births** ( $t = -0.08$ ) - about two orders of magnitude smaller than the naive (unadjusted) relationship. In other words, **the raw correlation between health spending and child mortality is essentially a wealth effect**: in this cross-sectional data, no independent effect of spending can be statistically detected once national wealth is accounted for. The direction (more spending, less mortality) was consistent across every check; the evidence for any causal magnitude was not.

## 1. The question and the data

Countries that spend more on healthcare tend to have lower child mortality. Does that mean spending *causes* lower mortality - or are both just symptoms of being a richer country? This is a genuine, unsettled question in health economics, and a clean example of confounding: GDP per capita plausibly drives both how much a country can spend on health *and* its mortality rate through completely different channels (nutrition, sanitation, education, water access). A correlation between spending and mortality is expected either way; a causal claim requires explicitly modelling and adjusting for that shared cause.

Data: Our World in Data's merged World Bank series - child mortality (WDI SH.DYN.MORT), current health expenditure per capita (SH.XPD.CHEX.PC.CD / .PP.CD, PPP-adjusted), and GDP per capita (NY.GDP.PCAP.CD). 186 countries with complete records for 2019 (aggregate regions and income-group rows excluded).

### Data preparation (Causality Agent, stage 1-2)

The agent's profiler flagged **country** and **country\_code** as identifiers (186 unique values in 186 rows - no causal signal) and **year** as constant within this single-year snapshot; both were accepted for removal. It also flagged outlier-heavy tails in **population** and **health\_expenditure\_per\_capita** (a small number of countries with extreme values), which were reviewed but left uncapped for this run.

| Column       | Action      | Why                            |
|--------------|-------------|--------------------------------|
| country      | drop column | Dropped column 'country'.      |
| country_code | drop column | Dropped column 'country_code'. |
| year         | drop column | Dropped column 'year'.         |

Result: 8 columns reduced to 5; 186 rows retained.

## 2. Proposed causal model (Causality Agent, stage 3)


With the cleaned dataset, the agent's language-model draft proposed the following structure, which was then statistically verified against the data (marginal correlation, partial correlation controlling for other variables, and pairwise data coverage):

*Agent's reasoning:* The proposed causal structure suggests that health expenditure per capita is a direct intervention that can lead to a reduction in child mortality rates, making it a plausible treatment candidate. GDP per capita serves as a confounding variable that influences both health expenditure and child mortality, indicating that wealthier nations may have better health outcomes due to increased spending on health. The region variable is included as it can impact both GDP and health expenditure, reflecting broader socioeconomic trends. The biggest structural uncertainty lies in the potential unobserved confounders such as healthcare quality and socioeconomic status, which may also significantly affect child mortality rates but are not measured in this dataset.

| Edge                                                     | Verifier verdict | Confidence | LLM's stated reason                                                                                                            |
|----------------------------------------------------------|------------------|------------|--------------------------------------------------------------------------------------------------------------------------------|
| health_expenditure_per_capita<br>-> child_mortality_rate | conflict         | 0.79       | Health expenditure per capita is expected to directly reduce child mortality rates by improving healthcare access and quality. |

| Edge                                            | Verifier verdict | Confidence | LLM's stated reason                                                                                                                  |
|-------------------------------------------------|------------------|------------|--------------------------------------------------------------------------------------------------------------------------------------|
| gdp_per_capita -> health_expenditure_per_capita | supported        | 1.00       | GDP per capita influences health expenditure per capita as wealthier countries tend to allocate more resources to health.            |
| gdp_per_capita -> child_mortality_rate          | supported        | 1.00       | Higher GDP per capita is associated with lower child mortality rates due to better overall living conditions and healthcare systems. |
| region -> gdp_per_capita                        | supported        | 0.83       | The region can affect GDP per capita as different regions have varying levels of economic development.                               |
| region -> health_expenditure_per_capita         | supported        | 1.00       | Different regions may have different health expenditure per capita due to varying government policies and economic conditions.       |
| region -> child_mortality_rate                  | supported        | 1.00       | Child mortality rates can vary significantly by region due to differences in healthcare access and economic factors.                 |

Note the nuance already visible here: the statistical verifier flagged the direct **health\_expenditure\_per\_capita -> child\_mortality\_rate** edge itself as “**conflict**” (confidence 0.79) even though the language model proposed it with high prior confidence - its raw partial correlation with mortality, after controlling for a few other variables, was weaker and less consistent than the correlations backing the other edges. That is an early warning sign for what the estimation and robustness sections confirm: the direct effect is real in direction, but not strong or stable in size.



Rendered causal graph as saved to the canvas.

### 3. Identification: what must be adjusted for

With the model above, the effect is identifiable via the **backdoor** criterion, requiring adjustment for: **region**, **gdp\_per\_capita**. Both treatment and outcome vary sufficiently in the data (no overlap violations), and the graph is a valid DAG (186 countries).

| Check               | Status | Detail                                          |
|---------------------|--------|-------------------------------------------------|
| Graph is a DAG      | pass   | Identification assumes an acyclic causal graph. |
| Treatment variation | pass   | Treatment must vary to support estimation.      |

| Check                | Status | Detail                                                     |
|----------------------|--------|------------------------------------------------------------|
| Outcome variation    | pass   | Low outcome variation weakens the analysis.                |
| Backdoor path review | warn   | Open undirected paths may indicate unresolved confounding. |

The “backdoor path review: warn” line refers to the direct treatment->outcome edge itself being listed as an open path, which is expected and correct (it must stay open - that is the effect being estimated); it is not evidence of unresolved confounding on top of the adjustment set above.

## 4. The estimate - and why adjustment matters here

Two versions of the estimate were computed to make the confounding concrete:

| Model                                     | Adjusts for                                                                       | Effect per \$1 of spending | Effect per \$1,000 of spending |
|-------------------------------------------|-----------------------------------------------------------------------------------|----------------------------|--------------------------------|
| Full model (adjusted)                     | region, gdp_per_capita                                                            | -1.54e-05                  | -0.0154                        |
| Direct-effects-only model (no adjustment) | none - caution: Identification exists, but unresolved backdoor paths remain open. | -1.61e-03                  | -1.6135                        |

The unadjusted model - which omits GDP as a cause of treatment, so no backdoor path is detected and nothing is controlled for - produces an effect roughly **104x larger in magnitude** than the GDP-and-region-adjusted estimate. That gap *is* the confounding: most of the naive correlation between spending and mortality is GDP acting through both variables, not spending causing mortality to fall directly. This is exactly the failure mode identification analysis exists to catch, and exactly why the adjustment set matters more than the point estimate on its own.

## 5. Robustness and refutation checks

Overall robustness score: **100%** (baseline method: backdoor.linear\_regression, baseline effect -1.54e-05).

### Refutation tests

| Test                | Result                      | Interpretation                                                                                                                                                                                |
|---------------------|-----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| placebo treatment   | status: ok, delta: 1.54e-05 | Effect collapsed to ~0 under a fake treatment, as it should - passes.                                                                                                                         |
| dummy outcome       | status: ok, delta: 1.38e-05 | Against a synthetic outcome the effect goes to ~0 (p=0.96) - passes.                                                                                                                          |
| random common cause | status: ok, delta: 2.40e-07 | Adding an irrelevant random variable barely moved the estimate - passes; the model is not sensitive to arbitrary extra noise.                                                                 |
| data subsample      | status: ok, delta: 2.76e-06 | On an 80% random subsample the estimate kept its sign in this run, but an earlier run's subsample flipped it - with an effect this close to zero, which countries land in the sample matters. |

### Estimator comparison

Propensity-score matching and weighting could not be run: both require a binary treatment, and health expenditure per capita is continuous. Linear regression is the appropriate estimator here and was used throughout.

## Confounder-strength sensitivity sweep

| Assumed unobserved-confounder strength | Adjusted effect |
|----------------------------------------|-----------------|
| 0.1                                    | -1.490e-05      |
| 0.2                                    | -1.436e-05      |
| 0.3                                    | -1.382e-05      |
| 0.4                                    | -1.328e-05      |
| 0.5                                    | -1.274e-05      |

The effect shrinks steadily toward zero as the assumed strength of an *additional, unmeasured* confounder increases. The formal partial- $R^2$  (Cinelli-Hazlett) sensitivity analysis sharpens this considerably: the adjusted effect's **robustness value is 0.0058**, meaning a confounder explaining just **0.58% of the residual variance** of both treatment and outcome would be enough to bring the estimate to zero - and the adjusted coefficient is not statistically significant to begin with ( $t = -0.079$ ). A factor as plausible as governance quality or health-system efficiency could easily account for the entire remaining effect.

## Competing-model comparison

Verdict: **moderate**. All 2 model variants agree on the direction of the effect (-0.0000, -0.0016), but the magnitude varies by 99%. The sign of the conclusion is trustworthy; quote the size of the effect with a caveat about model dependence.

## 6. Conclusions

- **Direction:** every check in this analysis - the adjusted estimate, the unadjusted estimate, the refuters, and the confounder sweep - agrees that higher health spending is associated with lower child mortality, never the reverse.
- **Magnitude:** the honest, GDP-and-region-adjusted point estimate is small (-0.015 deaths per 1,000 live births per \$1,000 of extra per-capita spending), roughly 100x smaller than the naive correlation, and not statistically distinguishable from zero.
- **Confidence:** low. The adjusted effect is not statistically significant ( $t = -0.08$ ), the formal sensitivity analysis shows a confounder explaining only 0.58% of residual variance would erase it entirely, repeated subsample refutations flip its sign between runs, and the competing-model comparison shows a 99% magnitude swing across plausible graph structures.
- **Policy read:** in this cross-sectional country-level data, once national wealth and region are accounted for, there is **no statistically detectable independent effect of health spending on child mortality**. That is itself the finding: the strong raw correlation is essentially a wealth effect. Detecting a genuine spending effect - which may well exist - would require stronger designs: within-country changes over time (the companion panel file), lagged spending, or individual-level data.

## Limitations

- **Ecological (aggregate) data:** this is country-level, not patient-level, data. A country-level effect does not prove the same causal story holds for individual children or households.
- **Cross-sectional, single year:** this run used the 2019 snapshot. The companion panel file (2000-2021, with a 2-year-lagged spending column) is available in this folder to test whether the relationship is stable over time or driven by a same-year assumption that may not hold.
- **Region encoded as a single ordinal covariate:** the current pipeline adjusts for region via integer category codes in a linear model rather than one-hot dummy variables, which is a simplification that could bias the region-adjustment component of the estimate.
- **Robustness score reads optimistically:** the dashboard's 100% score counts how many checks *completed*, not how favourable they were - the completed sensitivity analysis is precisely what shows the effect to be fragile.

## A note on producing this report

While assembling this analysis, cross-checking numbers between endpoints surfaced a real bug: two code paths were silently collapsing the continuous mortality-rate outcome into a binary 0/1 indicator before estimating - a transformation meant for genuinely binary outcomes (e.g. churn yes/no) that produces a meaningless number when applied to a continuous rate (the 186-value mortality series became 185 ones and a single zero). Both paths (agent/compare\_models and the "Run Inference" estimation) have since been fixed and verified: all estimation routes now agree on the adjusted effect of  $-1.54e-05$ . The formal partial- $R^2$  sensitivity analysis, reported as unavailable in section 5, has also been repaired since (it required benchmarking the hypothetical confounder against an observed adjustment variable) and now runs on this dataset.