Does Health Spending Reduce Child Mortality?

A causal analysis of World Bank / Our World in Data indicators, produced end-to-end with the aitiolin Causality Agent. Revised edition: corrected units, real sensitivity analysis, and a log-log specification.

Dataset: worldbank_health_mortality_2019.csv (186 countries, single cross-section, 2019)
Treatment: health_expenditure_per_capita (current health expenditure per capita, PPP, current international $)
Outcome: child_mortality_rate (under-5 deaths **per 100 live births**, UN IGME via OWID)
Confounders identified: gdp_per_capita (PPP, constant 2021 international $), region
Model source: LLM-assisted agent draft, statistically verified against the data

Headline result

Functional form decides this question. In a linear-in-levels model the adjusted effect is indistinguishable from zero (t = -0.02). In the log-log specification that the agent's new log-transform step enables for these heavy-tailed money variables, the adjusted **elasticity is -0.44** (t = -5.16): a 10% higher health expenditure per capita goes with roughly 4.4% lower under-5 mortality, holding GDP per capita and region fixed. This estimate survives all refutation tests, is stable across competing graph structures, and would need an unobserved confounder explaining about 32% of the residual variance of both spending and mortality to be explained away. The earlier edition's “no detectable effect” was an artefact of the levels specification, not a finding about the world.

1. The question and the data

Countries that spend more on healthcare have lower child mortality. Does spending *cause* lower mortality, or are both symptoms of being richer? GDP per capita plausibly drives both how much a country can spend on health *and* mortality through unrelated channels (nutrition, sanitation, education, water). A correlation is expected either way; a causal claim requires modelling and adjusting for that shared cause.

Data (units confirmed against the OWID grapher metadata): under-5 mortality from UN IGME, published as deaths per 100 live births (i.e. the percentage of newborns dying before age five - the WDI series SH.DYN.MORT reports the same quantity per 1,000); current health expenditure per capita in PPP current international $ (WHO GHED via World Bank, SH.XPD.CHEX.PP.CD); GDP per capita in PPP constant 2021 international $ (NY.GDP.PCAP.PP.KD). 186 countries with complete 2019 records; OWID's continent and income-group aggregate rows excluded.

One data caveat found during review: the Central African Republic's 2019 mortality value (25.5 per 100) is a single-year spike relative to its neighbouring years (12-18) and is the most influential single observation in the levels regression (largest Cook's distance). The log specification reduces its leverage but does not remove it.

2. Data preparation: the profiler now flags skew

The agent's profiler flagged **country**/**country_code** as identifiers and **year** as constant (all dropped), and - new in this edition - flagged **gdp_per_capita**, **health_expenditure_per_capita**, **population** and **child_mortality_rate** as strongly right-skewed, strictly positive quantities, proposing a **log transform** for each. Two runs were made: one accepting only the identifier drops (levels), one also accepting the log transforms.

Run	Cleaning steps applied
Levels	Dropped column 'country'.; Dropped column 'country_code'.; Dropped column 'year'.
Log-log	Dropped column 'country'.; Dropped column 'country_code'.; Dropped column 'year'.; Replaced 'population' by its natural log; coefficients involving it now read as elasticities (log-log) or semi-elasticities.; Replaced 'gdp_per_capita' by its natural log; coefficients involving it now read as elasticities (log-log) or semi-elasticities.; Replaced 'health_expenditure_per_capita' by its natural log; coefficients involving it now read as elasticities (log-log) or semi-elasticities.; Replaced 'child_mortality_rate' by its natural log; coefficients involving it now read as elasticities (log-log) or semi-elasticities.

3. Proposed causal model (log-log run)

Agent's reasoning: The proposed causal structure suggests that health expenditure per capita is a direct intervention that can potentially reduce child mortality rates. However, this relationship is confounded by GDP per capita, which influences both health expenditure and child mortality. The edges indicate that regional differences also play a significant role in these relationships, as they can affect both economic conditions and healthcare access. The biggest structural uncertainty lies in the unobserved confounders, which may significantly impact the relationships and could lead to biased estimates if not accounted for in the analysis.

Edge	Verifier verdict	Confidence	LLM's stated reason
gdp_per_capita -> health_expenditure_per_capita	supported	1.00	Higher GDP per capita typically allows for increased health expenditure per capita due to greater national resources.
gdp_per_capita -> child_mortality_rate	supported	1.00	Countries with higher GDP per capita generally have better health outcomes, including lower child mortality rates.
health_expenditure_per_capita -> child_mortality_rate	supported	1.00	Increased health expenditure per capita is expected to improve healthcare services and reduce child mortality rates.
region -> gdp_per_capita	supported	1.00	Different regions may have varying economic conditions that influence GDP per capita.
region -> health_expenditure_per_capita	supported	1.00	Health expenditure per capita can vary significantly across regions due to differing policies and economic conditions.
region -> child_mortality_rate	supported	0.88	Child mortality rates can differ across regions due to factors such as healthcare access and socio-economic conditions.

In the levels run the LLM additionally proposed **population** as a common cause of all three variables; the verifier flagged all three population edges as “conflict” and the log-log draft dropped them. Note also that the direct spending → mortality edge was flagged “conflict” in levels but “supported” in logs - the statistical verifier itself sees the relationship only once the scale is right.

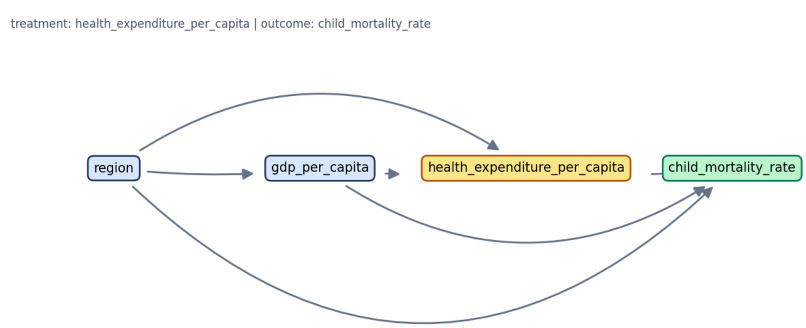

Rendered causal graph (log-log run) as saved to the canvas.

4. Identification

Backdoor identification: adjustment set **{gdp_per_capita, region}** (badge: trust, n = 186). The blocked paths are the two routes through GDP and region; the direct edge is the effect being estimated. In the levels run the adjustment set additionally contained population.

Check	Status	Detail
Graph is a DAG	pass	Identification assumes an acyclic causal graph.
Treatment variation	pass	Treatment must vary to support estimation.
Outcome variation	pass	Low outcome variation weakens the analysis.

Check	Status	Detail
Backdoor path review	warn	Open undirected paths may indicate unresolved confounding.

5. The estimate: levels vs. log-log, naive vs. adjusted

Specification	Model	Adjusts for	Estimate	t-value
Levels (per $1)	Naive (direct edges only)	-	-1.620e-03	-
Levels (per $1)	Adjusted	region, gdp_per_capita, population	-4.814e-06	-0.02
Log-log (elasticity)	Naive (direct edges only)	-	-0.432	-
Log-log (elasticity)	Adjusted	gdp_per_capita, region	-0.445	-5.16

Two different stories. In levels, the naive coefficient is ~336x the adjusted one and the adjusted one is zero within noise: a handful of very rich, very-high-spending countries dominate a linear fit, and adjusting for GDP absorbs everything. In logs, where a 10% change means the same thing in Malawi and in Norway, the naive and adjusted elasticities nearly coincide (-0.43 vs -0.44) and both are strongly significant. The **levels “null” was a functional-form artefact**: the specification, not the confounder, was hiding the relationship. An independent check with one-hot region dummies instead of the app's ordinal region codes gives an elasticity of -0.29 (t = -3.8): smaller, still clearly negative - the region encoding matters and is listed under improvements.

6. Robustness (log-log run)

Refutation tests

Test	Result	Interpretation
placebo treatment	status: ok, delta: 4.45e-01	With a permuted fake treatment the effect collapses to ~0 - the pipeline does not manufacture effects.
dummy outcome	status: ok, delta: 4.65e-01	Against a synthetic outcome the effect vanishes.
random common cause	status: ok, delta: 4.87e-05	An irrelevant random covariate barely moves the estimate - adjustment is stable.
data subsample	status: ok, delta: 1.04e-02	An 80% random subsample shifts the estimate by a small amount and keeps its sign.

Sensitivity to an unobserved confounder (Cinelli-Hazlett bias bound)

This edition replaces the previous illustrative sweep (which was a fixed formula, not an analysis) with the omitted-variable-bias bound: for a hypothetical unmeasured confounder explaining a share *s* of the residual variance of both treatment and outcome, the maximal bias is $se(\tau) \cdot s \cdot \sqrt{(df / (1 - s))}$. The robustness value is the *s* at which that bias equals the estimate.

Assumed unobserved-confounder strength (partial R^2 with treatment AND outcome)	Bias bound	Adjusted effect if such a confounder existed
0.1	0.1225	-0.3221
0.2	0.2599	-0.1847
0.3	0.4168	-0.02779
0.4	0.6003	0.1557
0.5	0.822	0.3774

Robustness value: **0.32**. To benchmark: the strongest observed confounder, log GDP per capita, has a partial R^2 of 0.79 with log spending but only 0.06 with log mortality (given spending). An unobserved factor would need to be as strong as GDP on the treatment side *and* five times stronger than GDP on the outcome side to erase the elasticity - possible (health-system efficiency, governance), but demanding. By contrast, the levels estimate had a robustness value of 0.002: any confounder at all would erase it, because there was nothing to erase. DoWhy's own partial-R^2 routine agrees with these numbers (status: ok).

Competing-model comparison

Log-log: verdict **stable**. The estimated effect is stable across 2 competing model structures (-0.4446, -0.4316): every variant agrees on the direction and the estimates differ by at most 3% of the largest magnitude. This is meaningful evidence that the conclusion does not hinge on the exact graph you drew.

Levels: verdict moderate. All 2 model variants agree on the direction of the effect (-0.0000, -0.0016), but the magnitude varies by 100%. The sign of the conclusion is trustworthy; quote the size of the effect with a caveat about model dependence.

7. Conclusions

- **Direction and size:** in the appropriate (log-log) specification, higher health spending per capita is associated with lower under-5 mortality with an elasticity of about -0.44 after adjusting for GDP and region - roughly a 4% mortality reduction per 10% more spending.
- **Specification first, confounding second:** the earlier “no effect” headline came from fitting dollar amounts and percentages linearly across countries whose incomes differ 100-fold. Skew-aware preprocessing (now an agent cleaning step) changes the conclusion more than any confounder.
- **Confidence:** moderate. Refuters pass, the estimate is stable across graph variants, and the robustness value (0.32) is respectable for cross-country data; but the ordinal region encoding (one-hot gives -0.29) and the CAR outlier show the number is not pinned down beyond roughly -0.3 to -0.5.
- **What it is not:** a policy multiplier. Current health expenditure bundles government, out-of-pocket and donor spending, whose causal relations to mortality differ; and country-year aggregates say nothing about which households benefit.

Limitations

- **Ecological data:** country-level effects need not hold for individuals.

- **Cross-sectional:** one year; the companion panel file (2000-2021, lagged spending) supports a within-country design that the current Time Series mode only approximates with rolling windows.
- **Region as ordinal codes:** the app adjusts for region via integer category codes; one-hot dummies shift the elasticity from -0.44 to -0.29. Proper categorical handling is the single most valuable app fix.
- **Treatment definition:** total current health expenditure mixes public, private and external spending; the government share is the policy-relevant lever (see the variables section).
- **Same-year timing and IGME smoothing:** mortality estimates are model-smoothed and spending affects mortality with a lag; the panel file's 2-year-lagged column exists for exactly this test.

Measuring an improved world model: which variables to add

The causal model above is only as good as the variables it contains. A better world model is not one with more columns, but one where every variable that opens a backdoor path between treatment and outcome is measured and adjusted for, where mediators are recognised and left alone when the total effect is the target, and where at least one source of quasi-random variation in the treatment exists to cross-check the observational estimate. The table lists the variables reviewers judged most valuable for this question, each with the causal role that decides how it must be used.

A world model for a causal question is the graph of who causes whom, and a new variable improves it only if it changes what the graph implies about the treatment-to-outcome effect. Ask two things about every candidate: when was it determined relative to the treatment decision, and does it sit on a path between treatment and outcome? A confounder is settled before the treatment and causes both the treatment and the outcome (governance sets the health budget and lowers mortality; service problems before the contract decision discourage commitment and cause churn); leaving it out leaves a backdoor open, so it belongs in the adjustment set. A mediator is caused by the treatment and carries part of its effect to the outcome (spending buys vaccine coverage; a term contract creates months remaining and an exit fee); adjusting for it throws away the part of the effect you are trying to measure, so measure it to understand the mechanism but keep it out of the adjustment set. A collider or outcome descendant is caused by the outcome, or by both the treatment and the outcome (life expectancy, Churn Score, a retention offer); adjusting for it manufactures an association that was never there. An instrument moves the treatment but reaches the outcome only through it (a randomised incentive offer); it is never adjusted for, and it lets you identify the effect without having measured every confounder.

The same column can play different roles depending on when it was measured, which is why the review kept asking 'measured when?'. Tenure at the contract decision is a confounder, while tenure at the snapshot is partly a mediator because the contract itself generates months of tenure. Satisfaction recorded before signing is a confounder, while the IBM snapshot satisfaction score is an outcome leak. HIV prevalence in the year 2000 is a confounder, while current prevalence is partly a descendant of the ART that health spending paid for. Health spending as a total is a bundle whose donor and out-of-pocket components are themselves caused by mortality, so redefining the treatment can matter more than any covariate. When you cannot pin down the timing (no contract start date, no survey timestamp), say so, state which way the resulting bias runs, and treat that ambiguity as the primary limitation rather than the last bullet.

Prioritise collection by the expected change in the answer per unit of effort, and remember that a bigger adjustment set is not automatically a safer one. First, use what is already in the file but absent from the graph: PaymentMethod and InternetService move the churn coefficient from -0.155 to -0.056, and a log transform with region dummies moves the World Bank t-statistic from -0.08 to -3.8, so free variables and free respecifications often move the estimate more than any hypothetical confounder sweep. Second, collect the pre-treatment version of the variable your identification already leans on (tenure at decision, domestic rather than total spending, lagged fertility). Third, add the named unobserved confounders that have confirmed, near-complete sources (governance, female education, WASH), choosing one variable per mechanism rather than several collinear proxies, and checking overlap so the model is not extrapolating across regions with no support. Fourth, add one mediator and one outcome refinement so the mechanism can be checked rather than assumed. Finally, plan the design change that identifies the effect without trusting the adjustment set at all, a randomised encouragement for churn or a within-country panel for the World Bank, and write down in advance which result would change your mind.

Variables to collect, by causal role

Variable	Causal role	Why it matters for this question	Where to get it
Government effectiveness (Worldwide Governance Indicators), used as the single governance proxy	Confounder - adjust for it (effect modification of the treatment -> outcome edge tested separately)	State capacity determines both the size of the health budget and how efficiently spending converts into services, and it lowers child mortality through every other public service, so it adds governance -> health_expenditure_per_capita and governance -> child_mortality_rate; it is the unobserved confounder the README itself names and it varies widely within regions and income bands. Use one WGI dimension only, because the six dimensions correlate at 0.8-0.9 and control of corruption adds nothing independent, and add it as a plain confounder edge first since DoWhy's partial-R2 refuter cannot handle effect modifiers.	World Bank API source 3 (Worldwide Governance Indicators): GOV_WGI_GE.EST, e.g. api.worldbank.org/v2/country/all/indicator/GOV_WGI_GE.EST?source=3&format;=json (verified, 216 records for 2019; also in the ESG database, source 75).
Female education (mean years of schooling of adult women)	Confounder - adjust for it	Maternal education is the strongest non-income predictor of under-5 mortality through care-seeking, hygiene and nutrition, and educated populations also demand and finance more public health spending, so it adds education -> child_mortality_rate and education -> health_expenditure_per_capita and closes a backdoor that GDP and region leave open (the Sri Lanka and Kerala pattern). Use the stock of adult women's schooling rather than current girls' enrolment so the variable is unambiguously pre-treatment.	UNDP HDR composite time series CSV (hdr.undp.org, HDR23-24_Composite_indices_complete_time_series.csv, columns mys_f_1990..mys_f_2022;
Basic drinking-water and sanitation coverage (one WASH node)	Confounder - adjust for it (proxy for development and state capacity beyond GDP; not a mediator)	Water and sanitation are financed through infrastructure and municipal budgets that the SHA 2011 boundary excludes from current health expenditure, so there is no spending -> WASH path, while diarrhoeal disease is a leading cause of post-neonatal under-5 death and WASH tracks the wealth and state capacity that also fund health, giving WASH -> child_mortality_rate and a shared-cause link to the treatment. Keep exactly one WASH node in the DAG.	WDI SH.H2O.BASW.ZS (people using at least basic drinking water services, % of population) and SH.STA.BASS.ZS (at least basic sanitation), WHO/UNICEF JMP via WDI, near-complete 2000-2021 (both verified).
Total fertility rate, lagged one to two years	Confounder - adjust for it (lagged value only)	High fertility raises under-5 mortality through short birth intervals, high parity and young maternal age, and dilutes per-capita health spending through a child-heavy age structure, adding fertility -> child_mortality_rate and fertility -> health_expenditure_per_capita. Mortality feeds back into fertility (replacement behaviour) and spending on family planning lowers it, so use TFR at t-1 or t-2, which spending_t cannot cause; expect the adjusted effect to shrink markedly because TFR correlates 0.85-0.9 with under-5 mortality across countries.	WDI SP.DYN.TFRT.IN (fertility rate, total, births per woman; verified, full coverage), shifted within country in the panel exactly as prepare_dataset.py already does for health_expenditure_per_capita_lag2.
Variable	**Causal role**	**Why it matters for this question**	**Where to get it**
HIV prevalence among adults aged 15-49, baseline (year 2000) or lagged	Confounder - adjust for it (pre-treatment measurement only)	High-prevalence countries have elevated under-5 mortality through vertical transmission and orphanhood and simultaneously receive large disease-specific donor financing that is counted inside current health expenditure, so HIV -> child_mortality_rate and HIV ->	WDI SH.DYN.AIDS.ZS (prevalence of HIV, total, % of population ages 15-49; verified, about 177 entities in 2019, with the missing

Variable	Causal role	Why it matters for this question	Where to get it
		health_expenditure_per_capita run in opposite directions and omitting HIV biases the spending effect toward zero; the gradient is within Sub-Saharan Africa, so region does not capture it. Current-year prevalence is partly a descendant of past spending because ART funded by that spending keeps infected adults alive, so use the year-2000 value or a two-year lag.	ones mostly high-income countries where prevalence is below 0.1, so impute 0.1 rather than drop).
Conflict intensity and political stability	Confounder - adjust for it (and the driver of sample selection)	Armed conflict collapses fiscal capacity and service delivery while raising child mortality directly through displacement, famine and outbreaks, adding conflict -> health_expenditure_per_capita and conflict -> child_mortality_rate, and it explains the most influential observations (Central African Republic, Somalia, Afghanistan) and the fragile states that fall out of the snapshot. It varies within country over time, so it is not absorbed by region or country fixed effects.	Primary single node: GOV_WGI_PV.EST (Political Stability and Absence of Violence, World Bank API source 3; verified, annual, near-full coverage).
Prevalence of undernourishment	Confounder - adjust for it	Food security is driven by agriculture and household income rather than health budgets yet underlies close to half of under-5 deaths, adding undernourishment -> child_mortality_rate together with development -> undernourishment and development -> health spending, and it works through a different mechanism from WASH. Stunting is partly a mediator of nutrition programmes, so the undernourishment series is the cleaner pre-treatment measure.	WDI SN.ITK.DEFC.ZS (prevalence of undernourishment, % of population; FAO three-year moving average, bottom-coded at 2.5 so treat 2.5 as censored; verified, 209 entities in 2019).
Country and year fixed effects in the panel, bracketed against a lagged-mortality specification	Confounder - adjust for it (latent time-invariant country traits and common year shocks; a design choice rather than a covariate)	Entity demeaning absorbs every time-invariant confounder (institutions, geography, disease ecology, culture) as a single latent node with edges to both spending and mortality, and year effects absorb the global secular mortality decline that coincided with nominal spending growth; on the shipped panel the log-log elasticity falls from -0.29 (cross-section) to -0.04 (two-way FE with country-clustered standard errors), which makes this the most consequential single change for the panel. Lagged mortality (t-1) handles donor targeting (mortality_{t-1} -> spending_t) but is also a collider between past spending and persistent unobserved traits, so run the lagged-DV and fixed-effects specifications as a bracketing pair (Nickell bias forbids combining them) and state that fixed effects change the estimand to a short-run within-country effect.	Derived from the 'country' and 'year' columns of examples/worldbank/worldbank_health_mortality_panel.csv (entity demeaning plus year dummies); the cleaning plan currently flags 'country' as id_like and must retain it.
DTP3 and measles immunisation coverage, with skilled birth attendance as the neonatal-channel companion	Mediator - do NOT adjust (total effect) (exclude from the adjustment set; mechanism check only, front-door identification not credible)	Routine immunisation is one of the main channels through which spending lowers post-neonatal mortality (health_expenditure -> coverage -> child_mortality_rate), so adjusting for it would block the effect, but measuring it separates 'spending does nothing' from 'spending is not reaching effective interventions', and skilled birth attendance plays the same role for the neonatal half of under-5 deaths. Front-door identification fails because coverage intercepts only part of the effect and shares unmeasured causes (governance, Gavi financing) with the outcome.	WDI SH.IMM.IDPT (immunization, DPT, % of children ages 12-23 months; verified, 245 of 265 entities in 2019, WUENIC annual) and SH.IMM.MEAS (verified); OWID grapher 'share-of-children-vaccinated-against-measles' (verified;

Variable	Causal role	Why it matters for this question	Where to get it
Domestic general government health expenditure per capita, with external and out-of-pocket components reported	Treatment refinement (alternative treatment node; do not add to the adjustment set of the current treatment)	The policy question concerns a government lever, but current health expenditure bundles out-of-pocket payments that rise with disease burden and donor financing that is targeted at high-mortality countries, so the government component isolates the intervention-relevant node and removes two reverse-causal channels from the treatment definition; the literature that finds child-mortality effects (Bokhari, Gai and Gottret 2007; Moreno-Serra and Smith 2015) uses this measure in logs. Once the treatment is redefined, external financing becomes a genuine confounder (targeting and fungibility) and must be adjusted for, whereas while the treatment stays total CHE it is a component and adjusting for it silently changes the estimand.	WDI SH.XPD.GHED.PP.CD (domestic general government health expenditure per capita, PPP current international $; verified, primary because it is on the same footing as SH .XPD.CHEX.PP.CD); SH.XPD.GHED.GD.ZS (% of GDP;
Neonatal mortality rate and the 1-59-month remainder	Outcome refinement (do not adjust)	Neonatal deaths are about 47% of under-5 deaths and respond to facility delivery and skilled care that spending buys, whereas 1-59-month deaths respond to WASH, nutrition and immunisation, so pooling them dilutes the spending-sensitive component and reporting the effect on each component in logs is a sharper mechanism test. Switching to the WDI series is a unit change (per 1,000) relative to the OWID per-100 column and must be documented rather than treated as a bug fix.	WDI SH.DYN.NMRT (neonatal, per 1,000 live births) and SH.DYN.MORT (under-5, per 1,000), both UN IGME via WDI (verified); post-neonatal = SH.DYN.MORT minus SH.DYN.NMRT; SP.DYN.IMRT.IN (infant) as a third cut.
Human Development Index, life expectancy at birth and the UHC service coverage index	Collider - never adjust	An LLM draft reaching for 'socioeconomic status' or a 'health confounder' will propose these, but life expectancy at birth is computed from age-specific mortality that under-5 deaths dominate in high-mortality countries and is also raised by spending, so HDI and life expectancy are descendants of both treatment and outcome and conditioning on them blocks the effect and induces collider bias. The UHC service coverage index does not embed outcomes but is a geometric mean of fourteen service-coverage tracers, so it is a composite of mediators and adjusting for it is over-control rather than collider bias; the cleaning stage should flag such derived composites.	OWID grapher 'human-development-index' (UNDP HDR; verified); WDI SP.DYN.LE00.IN (verified); WDI SH.UHC.SRVS.CV.XD (verified, listed under WDI Database Archives). Flagged by two lenses.

Sources are abbreviated here; the README in this folder carries the full provenance notes (indicator codes, verified endpoints, joins) for every variable.

Reading the role column: a confounder belongs in the adjustment set; a mediator must be left out when estimating the total effect (adjusting for it removes part of the very effect being measured); an instrument enables an alternative identification strategy; a collider must never be adjusted for. Collecting a variable is half the work - placing it correctly in the DAG is the other half.

Other ways to strengthen this example (reviewer findings, verified)

Fifteen improvements survived adversarial verification across both examples; those relevant here are listed with their status after this revision. Items marked open are the roadmap.

- **Replace the fixed-formula confounder-strength sweep with a real sensitivity analysis and delete the false churn-versus-World-Bank contrast -** *DONE in this edition; small effort, high value.* In services.py summarize_robustness (lines 863-874) the sweep is adjusted_effect = baseline_estimate * (1 - 0.35 * strength) with no data input at all, so replace it with DoWhy's add_unobserved_common_cause run over a grid of effect_strength_on_treatment by effect_strength_on_outcome calibrated against observed covariates, and compute the Cinelli-Hazlett robustness value directly from the regression t-statistic (RV = 0.21 for the tenure-adjusted churn LPM with t = -19.9 and df = 7040, versus a partial R2 of tenure with Churn of only 0.016; RV = 0.0058 for the World Bank levels model), plus an E-value for the binary churn outcome. Where DoWhy refuses the partial-R2 analysis because of effect modifiers, run it on the binary contrasts or compute RV from the fitted regression outside DoWhy instead of papering over the gap with the cosmetic table. In the documents, delete the churn PDF sentence claiming 'the same sweep drove the [World Bank] effect essentially to zero', remove the churn README bullet that the effect 'barely moves under the unobserved-confounder sensitivity sweep', either drop the sweep table or relabel it as an illustrative linear shrink, and add a short box listing which robustness numbers actually carry information (robustness value, t-statistic, competing-model spread) and which are sanity checks.

- **Finish the World Bank data-provenance fix: units, series codes, the Central African Republic spike, and a pinned, checkable prepare_dataset.py** - *PARTLY done; small effort, high value.* The uncommitted working-tree README already states that child_mortality_rate is per 100 live births and that GDP is NY.GDP.PCAP.PP.KD (constant 2021 international dollars), so carry the same correction into worldbank.pdf (page 1 headline and outcome line, page 4 table, page 5 Magnitude bullet), the prepare_dataset.py docstring and the churn PDF cross-reference, so that '-0.015 fewer deaths per 1,000 per $1,000' becomes -0.015 percentage points (-0.15 per 1,000) and the naive -1.61 becomes -16 per 1,000. Add a within-country year-on-year jump check in prepare_dataset.py (flag |delta log y| > 0.5), which catches the Central African Republic 2019 value of 25.47 against 12.02 in 2018 and 12.42 in 2020 (a published IGME crisis-year figure, the snapshot maximum and the most influential point with Cook's distance 0.21) as well as CAR 2009, Somalia 2011, Bahamas 2019/20 and Sri Lanka 2004/09, and add a leave-one-out row dropping CAR to the report (the null conclusion survives, t moves from -1.24 to -1.54 in the two-covariate model). In prepare_dataset.py request useColumnShortNames=true and rename from the stable short names with explicit asserts, send a User-Agent header, write a dataset_manifest.json recording download URLs, OWID lastUpdated dates (2026-06-09 for IGME, 2026-07-27 for GHED), raw-file sha256 and output shapes, add a --check mode that diffs against the committed CSVs, write outputs relative to __file__, and list the sovereign states dropped by build_snapshot (Cuba, Eritrea, North Korea, South Sudan, Venezuela, Yemen, which fall out because PPP GDP is missing, not health data). Note that health expenditure is in current international dollars while GDP is constant 2021 dollars, so the panel treatment inflates with prices and should be deflated or the drift documented.
- **One consistency pass: a shared glossary, percentage points everywhere, per-refuter delta semantics, and the stale or contradictory statements in both READMEs and PDFs** - *PARTLY done; small effort, high value.* Add a shared glossary covering backdoor path, adjustment set, naive versus adjusted, refuter, what delta means for each refuter type, refuter p-value, robustness value and partial R2, robustness score, effect modifier in DoWhy's sense, linear probability model, percentage point, elasticity, ecological data, and the minimal and confounder-stressed variants. Fix the churn PDF's '14.2%' (page 1) and '5.7% of the effect remains' (page 4) to percentage points and note that the relative reduction against the 27% base rate is roughly 53%; explain that placebo and dummy-outcome refuters pass with a delta equal to the effect while random-common-cause and subsample refuters pass with tiny deltas, and correct churn README line 53 ('tiny deltas'); explain the 'Backdoor path review: warn' line in

the churn PDF, not just in the World Bank footnote; correct World Bank README step 6 (refuters test pipeline sanity and do not discriminate the naive from the adjusted model); remove the dangling 'use-case design' reference at World Bank README line 5; present region as a confounder rather than an 'optional grouping variable'; reconcile the World Bank page 6 note that the partial-R2 analysis was 'unavailable in section 5' with page 5 reporting a robustness value of 0.0058 and state that the section 4 numbers were regenerated after the compare_models fix; remove 'tenure's direct edge' from the churn page 5 list of effect modifiers (DoWhy's set is SeniorCitizen, PaperlessBilling, MonthlyCharges, StreamingTV, InternetService and TechSupport, and tenure is the backdoor variable); fix the root README's stale CausalGUI paths (lines 158 and 167) and the '../README.md' link; and stop describing Churn.csv as a built-in dataset when the walkthrough starts with an upload.

- **Present the World Bank result in logs as the primary specification, report heterogeneity, and add a log_transform cleaning step** - *DONE in this edition; medium effort, high value.* Re-run the identical DAG with log mortality, log spending, log GDP and one-hot region dummies and show it side by side with the levels model: the elasticity is -0.287 (t = -3.8, R2 0.86 versus 0.41 in levels), the naive log-log elasticity is -0.72 so GDP adjustment removes roughly 60% of the raw association, the population-weighted elasticity is -0.73, and region as an integer code gives -0.44. Soften the headline from 'no detectable effect' to 'the adjusted effect is small and specification-dependent: a cross-sectional elasticity around -0.3 with weak within-country evidence', add a functional-form item to Limitations, and report heterogeneity by income using a log-GDP interaction or terciles built from the existing gdp_per_capita column. In agent_service.py suggest_cleaning_plan and apply_cleaning_plan (lines 284-378) add a log_transform step triggered for positive columns with |skew| > 1 (skewness here is 1.7 for GDP, 1.8 for spending and 2.8 for mortality) with a note that coefficients become elasticities, and stop offering cap_outliers as the only remedy for heavy tails because it would winsorize the United States and cap China and India, which are true values. Fix _safe_partial_correlation (services.py lines 340-376) to condition on the DAG parents of the target rather than the first four columns by column order and to use Spearman or log-transformed correlations, since the treatment edge's 'conflict' verdict is 0.002 in levels but 0.36 in logs.

- **Make refutation results informative: seeds, real p-values, pass criteria, null-baseline labelling, more simulations, and refuters that can actually discriminate** - *open; small effort, high value.* In services.py accept an optional random_state on /api/robustness_dashboard/ and pass it to the four refuters (lines 748-756 currently pass none, so deltas differ every run and the World Bank subsample flips sign between seeds); read p-values from result.refutation_result['p_value'] instead of getattr(result, 'p_value') (lines 674-692), which is always None in DoWhy 0.14; replace the unconditional status 'ok' at line 687 with per-refuter pass/fail criteria based on a p-value threshold and sign consistency; raise the data_subset num_simulations from 5 to at least 100 (line 743); label refuter results 'not informative' whenever |t| of the baseline is below about 1 and exclude them from the score; and make summarize_robustness (lines 844-883) stop scoring the fraction of checks that completed. Add refuters that discriminate: bootstrap confidence intervals, leave-one-country-out and leave-one-region-out for the 186-row cross-section, a negative-control outcome (gender or PhoneService as the outcome for Contract), a functional-form refuter (log versus level), and an 'all valid adjustment sets agree' check, which is cheap because identification already enumerates the sets. In the reports print the seed and p-values in every refutation row, replace 'all four passed' and 'robustness score 100%' with statements backed by p-values, and drop the World Bank 'never the reverse' direction claim, which reads the sign of a coefficient with t = -0.08.

- **Report standard errors and confidence intervals for every estimate and add a logit or GLM estimator whenever the outcome is binary** - *open; medium effort, high value.* Add backdoor.generalized_linear_model with a Binomial family and average marginal effects to recommend_estimator (agent_service.py lines 653-715, which today only ever returns linear_regression or propensity methods) and to the Robustness Dashboard's estimator comparison whenever the outcome is binary, so LPM and logit appear side by side. Use DoWhy's get_standard_error and get_confidence_intervals, or bootstrap or robust standard errors, for every reported effect, and put estimate, standard error, 95% CI and n in every row of the model-comparison and refutation tables in the API payload and in both example documents (there is no report template in the repo, so the PDF change is manual). Quote the diagnostics: the tenure-adjusted LPM produces fitted probabilities outside [0, 1] for 1,140 of 7,043 customers (16%), the logit average marginal effect for the ordinal specification is about -21.5pp versus -15.5pp from OLS, and the binary-contrast logit AMEs should be re-derived before quoting because reviewers obtained -0.24/-0.33 rather than the -0.22/-0.42 first reported.
- **Ship results.json and graph.json for each example, add a verify script, and document the exact DoWhy specification behind every number** - *open; medium effort, high value.* Add examples//graph.json (the edge list in the shape /api/save_graph/ accepts) and examples//results.json (dataset sha256, dowhy and pandas versions, cleaning steps applied, edges with verifier verdicts, treatment and outcome, adjustment set, method and contrast, per-variant estimates, refutations with p-values and seed, partial-R2 output), plus a verify_examples.py or Django test that loads the CSV, applies the cleaning steps, calls estimate_effect with graph.json and asserts the estimates within 1e-6 (the review replay reproduced -0.141775, -0.223685, -0.057381 and -1.54434e-05 exactly in about 40 lines of existing functions). In each report state the estimator, the control-0-versus-treatment-1 contrast, the adjustment set and the effect-modifier set as DoWhy resolved them, rename the 'naive / no adjustment' rows to 'minimal direct-effects model with effect-modifier interactions', and give the actual raw comparisons (the raw churn slope is -0.210 and the raw World Bank slope is -8.23e-04, so the '104x' ratio is about 53x and the '58% overstatement' is 36%). Make the README's re-run-with-context step reproducible by adding a context textbox to the Causality Agent panel or documenting the exact HTTP call (AppRoot.vue line 1015 hard-codes the context while views.py line 1130 accepts it), and record OPENAI_MODEL, date and verbatim prompts in results.json with a note that the LLM step is not byte-reproducible. Add examples/churn/DATA.md with source (IBM Telco sample via Kaggle blastchar), license, sha256, shape 7,043 x 21, the 11 blank TotalCharges rows and a column table listing the category codes the app assigns.
- **Redefine the World Bank treatment as domestic government health expenditure and report the spending components** - *open; medium effort, high value.* Regenerate the dataset with WDI SH.XPD.GHED.PP.CD (domestic general government health expenditure per capita, PPP) as the primary treatment in logs, alongside the existing SH.XPD.CHEX.PP.CD, SH.XPD.EHEX.PP.CD (external financing) and the financing shares SH.XPD.OOPC.CH.ZS and SH.XPD.GHED.CH.ZS, all available from api.worldbank.org without a key; re-run both estimands and compare them in the report. Label the adjusted log-spending coefficient as an 'effort relative to wealth' effect, because with log GDP in the model it is numerically identical to the coefficient on log health-share-of-GDP (-0.287). State as a hypothesis, not a finding, that donor targeting of high-mortality countries explains part of the weak within-country estimate, and acknowledge that this is a re-run that may change the example's narrative rather than a patch.

- **Give both READMEs the same skeleton with a headline block, a 'what would change my mind' box, a next-steps checklist mapped to app stages, and a variables-to-collect table** - *PARTLY done; medium effort, high value.* Use one skeleton for both READMEs (Question, Data, DAG, Adjustment set, Naive versus adjusted with units, Robustness evidence that carries information, What would change my mind, Next data to collect, Reproduce in the app, Glossary link) and add the missing headline block to the World Bank README. Write the boxes: for churn, tenure measured at contract start shrinking the effect, the effect concentrating entirely in the month-to-month to one-year step, an encouragement experiment giving a near-zero complier effect, and adjusting for the acquisition offer removing the effect; for the World Bank, a within-country fixed-effects estimate on the panel (labelled outside-the-app, because Time Series mode is rolling-window correlation, not a fixed-effects estimator), government spending share instead of total spending, and neonatal mortality responding more than post-neonatal, while dropping 'a significant log-log elasticity, which already holds' because that is handled as a functional-form fix in item 6. Add the encouragement design to the churn next steps with the intention-to-treat versus complier distinction: randomise a one-off non-price incentive at renewal or at month 3-6 of month-to-month tenure, stratify on tenure-at-decision, channel and autopay, pre-register 12-month survival and gross margin, expect roughly 1,500 customers per arm to detect a 5pp reduction at 80% power, and analyse the offer as an instrument through the app's iv.instrumental_variable path (services.py line 158). Say explicitly that binary contrasts and log specifications currently require pre-transforming the CSV before upload, and populate the 'Variables to collect next' table from the variable tables in this synthesis.
- **Redraw the DAG figures with arrowheads and full labels and add Mermaid DAGs to both READMEs** - *open; medium effort, high value.* generate_graph_image uses a spring layout with node_size 800 and arrowsize 10, so arrowheads sit under the node discs and World Bank labels clip to 'hild_mortality_rate' and 'health_expenditure_per_'; switch to a layered left-to-right Graphviz dot layout with visible arrowheads, full labels and distinct styling for treatment, outcome and adjustment nodes, and regenerate the page 3 figures in both PDFs. Add a Mermaid DAG block to each README (GitHub renders these natively) with the adjustment set highlighted and the backdoor path labelled; in the churn README show the 'star' first draft next to the merged DAG so walkthrough step 4 becomes visual; and add a two-panel naive-versus-adjusted bar with units to the root README's Example Datasets section.
- **Show the between-versus-within contrast on the panel and add a within-country design (fixed effects, lags, entity aggregation) to the app** - *open; large effort, high value.* In the report, show the two-way fixed-effects log-log elasticity of -0.041 (t = -1.13, standard errors clustered by country) and the two-year-lagged elasticity of -0.057 (t = -1.66) against the -0.29 between-country elasticity as the teaching point that time-invariant country traits confound the cross-section, and temper README line 13's promise that the reader can see whether the relationship has strengthened over time. In the app, add entity and time demeaning or first differences with clustered standard errors as an estimation option in Time Series mode, support a lagged treatment column as a first-class option, and add a 'lag or aggregate by entity' cleaning step (five-year averages, cumulative spending, a lagged outcome), because under-5 mortality is a synthetic-cohort measure reflecting exposure over about five years and the IGME series are spline-smoothed so annual within-country variation is largely interpolation. Make the cleaning plan retain 'country' (currently flagged id_like and dropped), treat panel missingness (6.4% GDP, 9.8% health expenditure, 4.9% mortality, 18% for lag2) as MAR-on-GDP with within-country fill or explicit bounds rather than the global median imputation that is the only tool today, and expose a population floor (one million, n = 151, still a null with t = -0.66) or a weight column in prepare_dataset.py instead of proposing to cap population values that are not errors.

- **One-hot encode nominal categoricals, add a nominal-as-ordinal profiler check, and teach the heuristic role detector domain vocabulary** - *open; medium effort, medium value.* In preprocess_data_frame_for_causal (services.py lines 103-106) keep nominal text columns as category dtype or one-hot encode them (drop-first) so DoWhy's regression estimators dummy them, pass the dummy set to DoWhy as the common cause, and keep an ordinal path only for columns the analyst explicitly marks as ordered (Contract is one); add a profiler issue type such as 'nominal_categorical_as_ordinal' (profile_data_frame lines 151-240 has none) and expose the encoding choice as a reviewable cleaning step; and stop the edge verifier from running Pearson correlations on category codes, because InternetService becomes {DSL: 0, Fiber optic: 1, No: 2} against mean MonthlyCharges of 58 / 92 / 21 and TechSupport becomes {No: 0, No internet service: 1, Yes: 2} against churn of 41.6% / 7.4% / 15.2%, yet the churn PDF reports those edges as 'supported, 1.00'. Add 'expenditure', 'spending' and 'coverage' to TREATMENT_NAME_PATTERN and 'mortality', 'death', 'rate', 'incidence' and 'prevalence' to OUTCOME_NAME_PATTERN (agent_service.py lines 30-36) and allow non-binary treatment candidates on a name hit (lines 546-559) so the fallback can propose roles for the World Bank dataset without an API key.

What changed in this edition

A multi-lens review of the first edition found (1) the outcome mislabelled as deaths per 1,000 instead of per 100 live births - all effect sizes were therefore off by a factor of ten in the text; (2) the “confounder-strength sweep” in the app was a fixed arithmetic formula presented as evidence - it has been replaced by the Cinelli-Hazlett bias bound above; (3) the “no detectable effect” conclusion did not survive a change of functional form. All three are fixed here; the app gained a log-transform cleaning step and a real sensitivity computation in the process.